

PCMCsim: An Accurate Phase-Change Memory Controller Simulator and its Performance Analysis

ISPASS 2022

Hyokeyun Lee*, Hyungsuk Kim*, Seokbo Shim[‡], Seungyong Lee*,
Dosun Hong[‡], Hyuk-Jae Lee*, Hyun Kim^P

* Seoul National University

^P Seoul National University of Science and Technology

[‡] SK Hynix



Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ Our simulator: PCMCsim
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

Outline

◆ Background & Goal

◆ Motivation

◆ Our simulator: PCMCsim

- System overview
- Execution flow
- Optimization: AIT Manager Prefetcher

◆ Evaluation


- Accuracy validation
- Takeaway messages

◆ Conclusion

Background


◆ Phase-change memory (PCM or PRAM)

- Non-volatile + High density
 - Material phase \rightarrow Resistance value
 - “1T1R” \rightarrow No capacitor
- Promising next-generation memory/storage



◆ Use of PCM

- Fast storage
 - Mount DAXFS
 - Product: Intel Optane DC SSD
- Low-cost main memory pool expansion
 - Further reduce cost \rightarrow disaggregated memory
 - Product: Intel Optane DC PMM



Simulation of PCM in previous studies

- ◆ Modifications on processor simulators (e.g., McSimA+)
 - Modify memory latency values
 - Few timing parameters → low simulation accuracy
- ◆ Modifications of DRAM simulators (e.g., DRAMsim)
 - More timing parameters, modified in detail
 - Slow (cycle-level) + No features for PCM controller
- ◆ Dedicated simulator (e.g., NVMain)
 - Simulation of cell-level characteristics
 - vs. modern PCM module? → Old structure...

Our goals

◆ Develop a PCM controller simulator achieving following three goals

◆ Goal1: Incorporating all necessary features of modern PCM modules

- (1) Hardware path; (2) address indirection table (AIT); (3) read-modify-write (RMW)
- Different from other memory controller architectures

◆ Goal2: Fast and accurate PCM controller simulator

- Event-driven is generally faster than cycle-level
- Accuracy validation against real hardware is required

◆ Goal3: Performance characterization of new simulator

- New architecture → new performance characteristics
- Providing key takeaway messages + Optimization

Outline

- ◆ Background & Goal
- ◆ **Motivation**
- ◆ Our simulator: PCMCsim
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

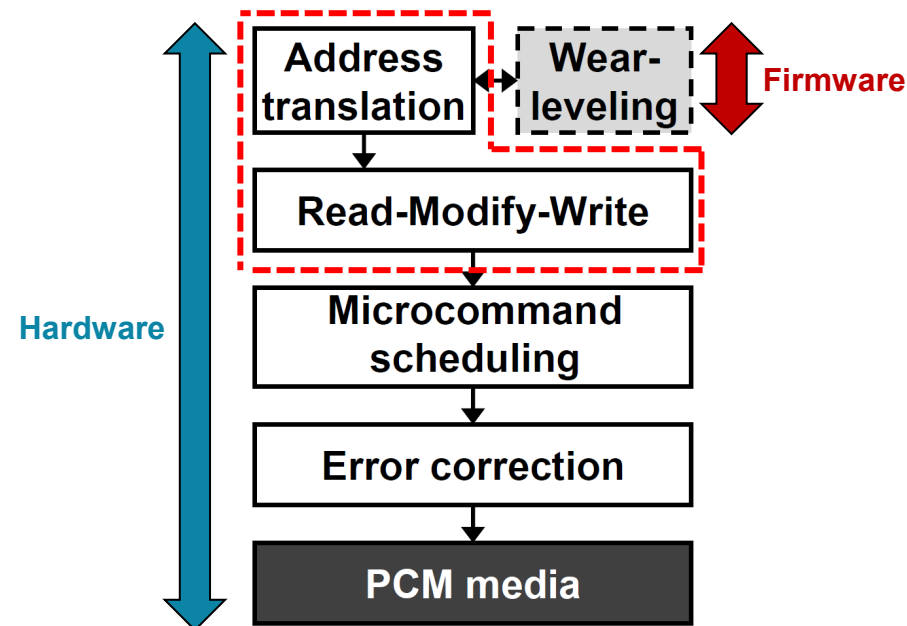
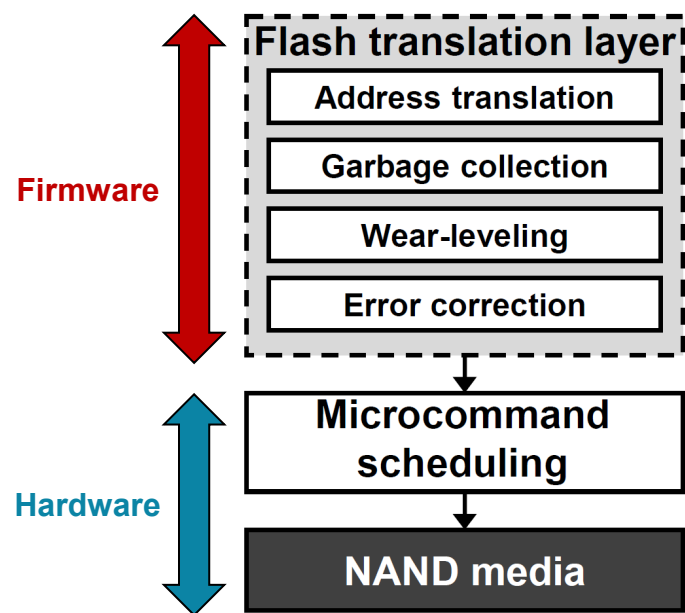
Necessity of PCM controller features

◆ Three features are required for a modern PCM controller

- AIT: mapping table that translates host address to device address
- RMW: manage gap between host cacheline size and PCM access granularity
- Hardware path: process host requests mainly with hardware pipeline

◆ Reasons for incorporating three features

- NAND mainly processes with firmware → wrong performance characteristics
- AIT and RMW account for large overhead → new research directions



Reducing latency of accessing AIT

- ◆ Accessing AIT incurs high latency overhead
- ◆ An address translation hardware needs to access DRAM-subsystem
 - More than 20% of total latency
 - Considerable in PCM-based systems
 - Naively thinking: **DRAM latency = 44.6 ns** vs. **PCM latency = 150 ns**
 - Practical: go through not only DRAM but also interconnect
 - Also, our evaluation result (later) shows similar results
- ◆ A low-cost speedup scheme for AIT access is required

Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ **Our simulator: PCMCsim**
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

Our simulator: PCMCsim

◆ PCMCsim?

- Simulates features for modern PCM controller architecture
- Event-driven simulator
- Functional- and cycle-accurate simulator

◆ Modular implementation for scalability

- Instantiation of a DRAM subsystem for AIT
- Capable of simulating DRAM in a standalone manner

◆ Facilitates design exploration of various capabilities of PCM

Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ **Our simulator: PCMCsim**
 - **System overview**
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

System overview

◆ Hardware request processing path (including RMW)

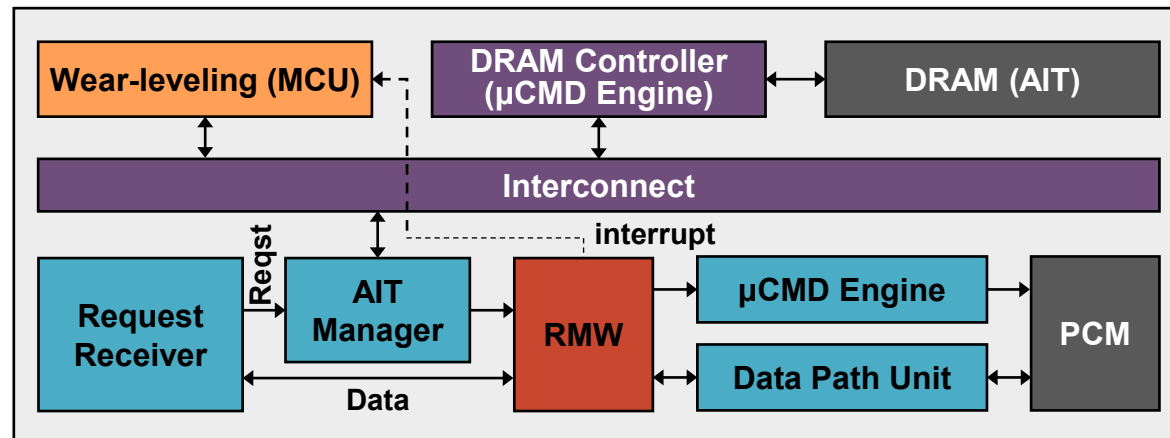
- Hardware pipeline for processing requests
- Major difference compared to NAND controller

◆ DRAM subsystem

- Path for accessing AIT in DRAM requested by **AIT Manager** or **wear-leveling**
- DRAM controller is simply instantiated using uCMD engine

◆ RMW

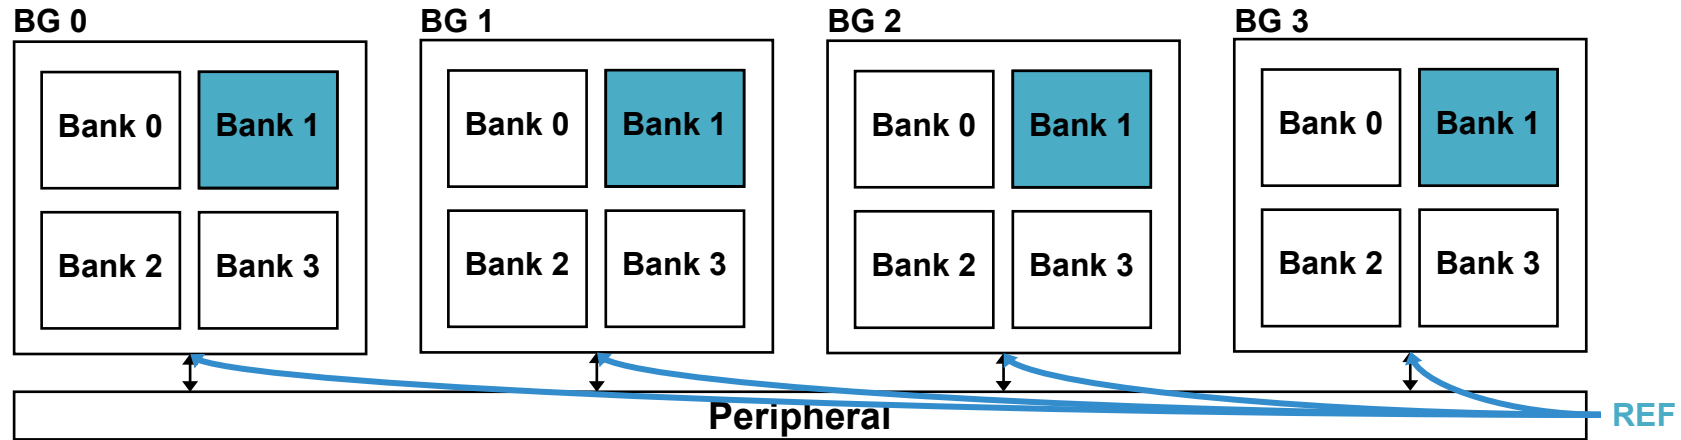
- Read data as PCM access granularity (≥ 128 B) + partial update (64B)
- Merge write data of same address / directly serve read request



More features in uCMD engine

◆ DDR5 features are supported

- **Same-bank refresh**: refresh banks with same index, across all bank groups



- Diverged timing parameter of **tCCD_L_WR**
 - **Def.:** Turn-around latency of two successive write commands in same bankgroup
 - tCCD_L_WR: RMW latency → if total prefetch length ≠ on-die ECC codeword length
 - tCCD_L_WR2: just-write (JW) latency → if total prefetch length == on-die ECC codeword length

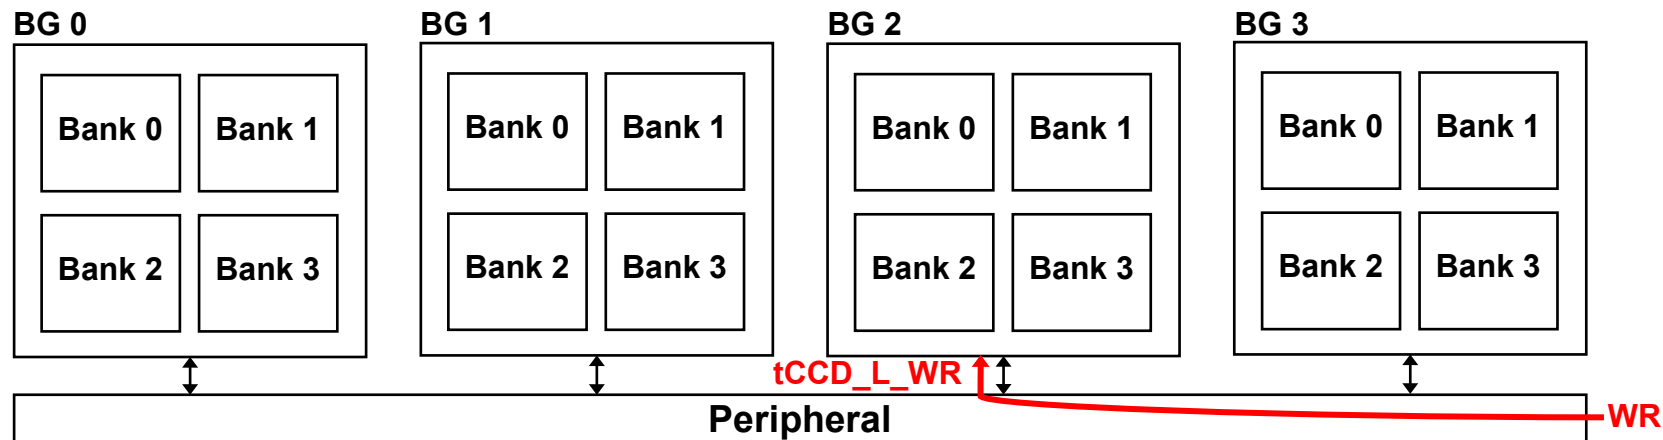
◆ Low-power features are supported

- Self-refresh, fast- and slow-exit powerdown modes
- Facilitate future low-power research in a DRAM system

More features in uCMD engine

◆ DDR5 features are supported

- **Same-bank refresh**: refresh banks with same index, across all bank groups



- Diverged timing parameter of $t_{CCD_L_WR}$
 - **Def.:** Turn-around latency of two successive write commands in same bankgroup
 - $t_{CCD_L_WR}$: RMW latency → if total prefetch length \neq on-die ECC codeword length
 - $t_{CCD_L_WR2}$: just-write (JW) latency → if total prefetch length $==$ on-die ECC codeword length

◆ Low-power features are supported

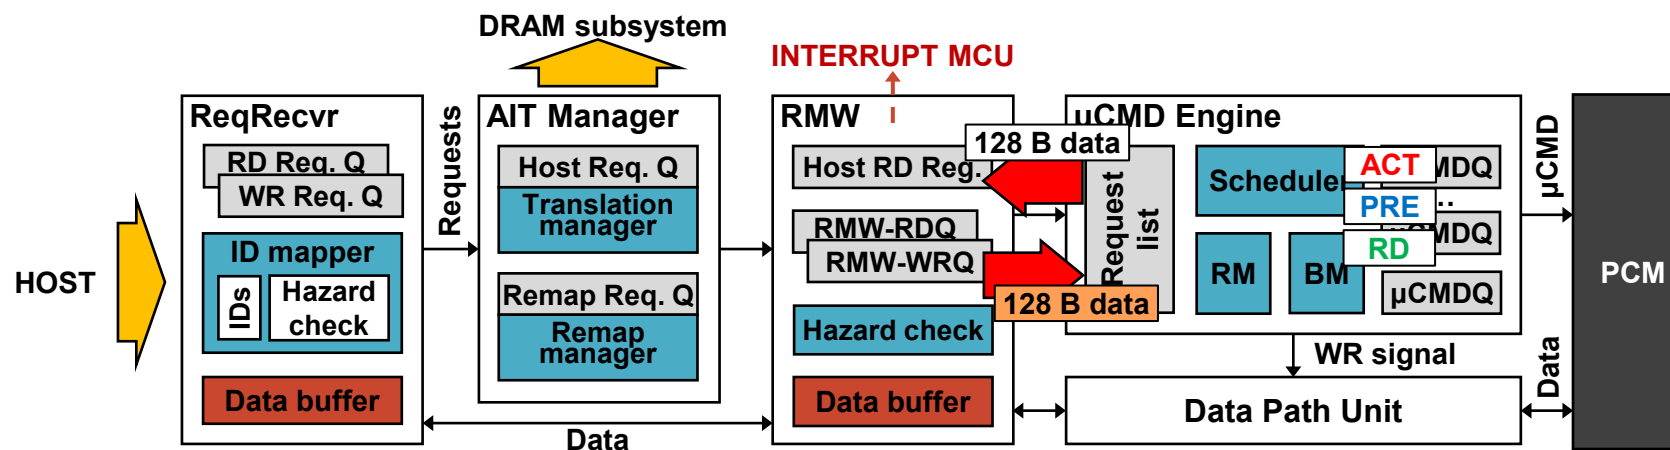
- Self-refresh, fast- and slow-exit powerdown modes
- Facilitate future low-power research in a DRAM system

Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ **Our simulator: PCMCsim**
 - System overview
 - **Execution flow**
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

Execution flow: hardware pipeline

- ◆ Request Receiver (ReqRecvr) stages host requests
 - Nearest interface for the host
 - Allocate request ID & Hazard handling (RAW, WAW, WAR)
- ◆ Two functionalities in AIT Manager
 - Function 1: address translation by accessing DRAM subsystem
 - Function 2: update AIT on DRAM subsystem & generate swap requests to PCM-side
- ◆ RMW
 - Reads data with PCM access granularity → partial update → write back
 - Interrupt MCU if write counts > THRESHOLD
- ◆ uCME engine → decompose **request** into **micro-commands**




Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ **Our simulator: PCMCsim**
 - System overview
 - Execution flow
 - **Optimization: AIT Manager Prefetcher**
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ Conclusion

Optimization: AIT Manager Prefetcher

- ◆ Accessing AIT in DRAM incurs significant performance overhead
 - Contributes more than **20% of total latency** (shown later)
- ◆ AIT Manager Prefetcher → reduce the access latency
 - Stream table → detects stream pattern
 - $\text{delta} = \text{input_address} - \text{current_address}$
 - $|\text{delta}| \leq \text{distance tolerance} \rightarrow \text{update current_address} + \text{pass previous entry value to PF}$
 - $|\text{delta}| > \text{distance tolerance} \rightarrow \text{replace the LRU entry with input_address}$
 - Prefetch filter → prevents generating same prefetch requests
 - $\text{valid} == 0$: part of new stream pattern → prefetch new stream pattern
 - $\text{valid} == 1$: prefetched previous “stream” → prefetch next stream pattern

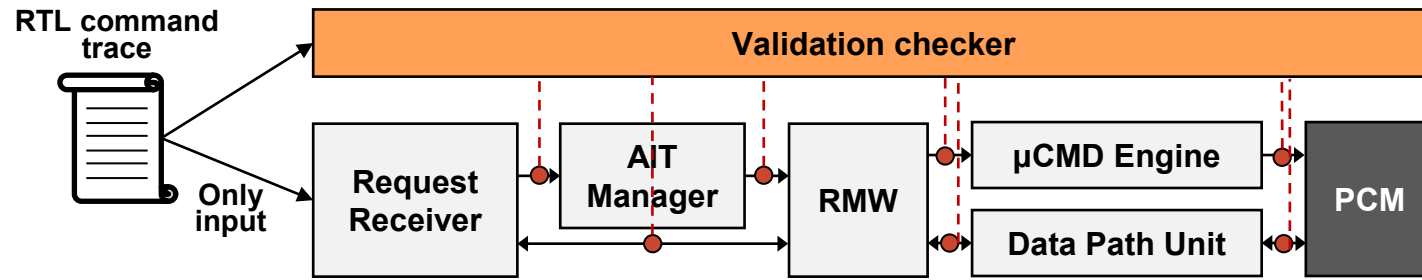


Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ Our simulator: PCMCsim
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ **Evaluation**
 - **Accuracy validation**
 - Takeaway messages
- ◆ Conclusion

Accuracy validation

- ◆ Validated cycle accuracy and functional accuracy
- ◆ Inject hardware RTL command trace to PCMCsim and validation checker
 - Info: > 2K commands with read/write ratio of 1.25 and various access patterns
 - Functional accuracy → probed commands between two modules
 - Cycle accuracy → compared simulated cycles and timestamps in RTL trace



- ◆ Functional error: 3.16%; cycle error: 13.6%

Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ Our simulator: PCMCsim
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ **Evaluation**
 - Accuracy validation
 - **Takeaway messages**
- ◆ Conclusion

Methodology

◆ Simulator

- gem5 (for memory trace extraction), PCMCsim (as a main memory)

◆ Workloads

- SPEC CPU 2017 memory trace (extracted from gem5)
- Umass OLTP Application I/O (UMASS) traces


◆ System parameters

- gem5: Intel i7-7700
- Baseline in PCMCsim (see below)

ReqRecv		single queue, 128 request IDs available
AITM	logic	2-entry command queue
	DRAM-AIT	DDR4-2666, 4 GB, $\times 16$ device, 64 B line 1 rank, 2 bank groups, 4 banks/bank group, tRAS=35 ns, tRCD=14.25 ns, tWR=15 ns, tCL=14.25 ns, tRP=14.25 ns,
RMW		2-entry data buffer for read-modify-write
uCMDE (PCM)		DDR4-2666, 128 GB, $\times 16$, 128 B line 1 rank, 2 bank groups, 4 banks/bank group, tRAS=100 ns, tRCD=100 ns, tWR=150 ns, tCL=14.25 ns, tRP=14.25 ns

Observation 1: latency breakdown


- ◆ Accessing AIT contributes 20% of total latency for both workloads
- ◆ SPEC CPU 2017 traces
 - No write requests → negligible RMW overhead
- ◆ UMASS
 - Many write requests → RMW contributes 25% of total latency



Message 1: architectural bottlenecks in a modern PCM controller newly come from AIT and RMW

Observation 2: Request Receiver


- ◆ Sweeping watermarks of read & write queues in Request Receiver
 - RxWy: issue WR after x RDs; issue RD after y WRs
 - Round Robin (R1W1): alternately issue read request and write request
- ◆ No latency variation for SPEC CPU 2017 → **No write requests**
- ◆ Latency increases (i.e., lower QoS) as the priority of read request lowers
- ◆ **No effects on throughput** → determined by WAWs in RMW than ReqRecvr



Message 2: watermark configuration in Request Receiver determines QoS rather than throughput for write-intensive workloads

Observation 3: AIT Manager


- ◆ Sweep host request queue size in AIT manager
- ◆ Larger queue → lower latency due to **reduced backpressure** on the pipeline
- ◆ More improvement variances for UMASS due to R/W ratio → more details in paper
- ◆ Increasing queue size amplifies the latency of modules



Message 3: host request queue size in AIT manager, which is cost-effective at 8, is a key parameter to reduce pipeline backpressure on Request Receiver

Observation 4: RMW

- ◆ Sweep data buffer size in RMW
- ◆ No latency variation for SPEC CPU 2017 → **No write requests**
- ◆ Larger data buffer → **lower latency**
- ◆ Some workloads → **small QoS improvement** due to relatively low merge ratio



Message 4: for write-intensive workloads (e.g., UMASS), increasing the number of data buffer entries in RMW significantly improves the performance because of merge operation in RMW


Observation 5: prefetcher performance

◆ Performance of different prefetcher schemes

- Previous studies: stream (i.e., next-line), SPP, VLDP
- AMP+AIT\$(small): AMP and 1 KB of AIT cache
- AMP+AIT\$(large): AMP and 64 KB of AIT cache

◆ Different prefetchers are suitable for different workloads

- AMP is better for SPEC CPU 2017 → highest prefetch accuracy
- AMP is worse for UMASS → lower prefetch accuracy for detecting stream patterns



Message 5: applying prefetcher schemes in AIT manager significantly improves QoS (up to 68.5%), but it is recommended to combine different prefetch schemes for general applications (i.e., no silver bullet)

Outline

- ◆ Background & Goal
- ◆ Motivation
- ◆ Our simulator: PCMCsim
 - System overview
 - Execution flow
 - Optimization: AIT Manager Prefetcher
- ◆ Evaluation
 - Accuracy validation
 - Takeaway messages
- ◆ **Conclusion**

Conclusion

- ◆ PCMCsim is a new simulator for modern PCM controller
 - AIT (including DRAM subsystem), RMW, hardware processing path are included
 - High simulation accuracy
 - Supports newest standard (DDR5) and low-power features for general-use
- ◆ Characterize the performance of modern PCM controllers
 - Provide insights for modern PCM controllers
 - AIT and RMW contribute majority of the latency

Thank you!



For information...

- Q&A / email: hklee@capp.snu.ac.kr
- More details in paper
- Code is available at: https://github.com/harrylee365/pcmcsim_public